

Instrumenting a Matching-Market Agent-Based Simulator for Broker-Advantage Attribution: Model Changes, Diagnostic Controls, and Results

Work in progress — progress report

April 19, 2026

Abstract

We instrument `TransientBrokerage.jl` — a base-model matching-market ABM — to measure when a central broker provides *pair-conditioned informational* value to agents, as distinct from a pure *data-pooling* or *capacity* advantage. Standard holdout metrics conflate the three. We add (i) a decoupled fee/cost split, (ii) per-period demand/fill/value counters with an algebraic decomposition into quality-selection, fill-access, and fee-cost components, (iii) pre-match access/assessment flags, (iv) seven ablation modes — including a `BLINDBROKER` that zeros the focal-agent input while preserving pooled training, (v) pooled and within-agent de-meaned holdout R^2 /rank metrics, and (vi) a two-way residualized holdout metric that isolates the pair-specific signal. Applied to diagnostic cells at $N=1000, T=200$, the instrumentation shows that the broker’s apparent advantage under a degenerate additively-separable DGP ($\rho=1$, "EasyQuality") is 94% a data-pooling artifact, while at the regime-gated anchors (Default, HardPairwise) pair-conditioned input drives most of the gap. The existing scripts remain byte-identical on the canonical seed.

1 Base model

Each period t a population of N agents with types $\mathbf{x}_i \in \mathbb{R}^d$ on the unit sphere demand matches. For a pair (i, j) the deterministic match signal is

$$f(\mathbf{x}_i, \mathbf{x}_j) = \underbrace{\rho \cdot \frac{1}{2} (\mathbf{x}_i^\top \mathbf{c} + \mathbf{x}_j^\top \mathbf{c})}_{\text{additive quality}} + \underbrace{(1 - \rho) \cdot g(\mathbf{x}_i, \mathbf{x}_j) \cdot \mathbf{x}_i^\top A \mathbf{x}_j}_{\text{gain-modulated interaction}}, \quad (1)$$

with the regime gain

$$g(\mathbf{x}_i, \mathbf{x}_j) = 1 + \delta \operatorname{sign}(\mathbf{x}_i^\top B \mathbf{x}_j). \quad (2)$$

A is SPD, B is symmetric and weighted-orthogonalised against A , $\mathbf{c}$ is an "ideal type" drawn from a sinusoidal curve, $\rho \in [0, 1]$ trades additive quality against pairwise interaction, and $\delta \in [0, 1]$ controls the magnitude of the two-regime gate. Realised output is $q_{ij} = Q + f(\mathbf{x}_i, \mathbf{x}_j) + \varepsilon_{ij}$ with $\varepsilon_{ij} \sim \mathcal{N}(0, \sigma_\varepsilon^2)$ and offset $Q = 1$. Calibration fixes $q_{\text{cal}} = \mathbb{E}[Q + f]$ by Monte Carlo and $r = 0.6 q_{\text{cal}}$ is the outside option. Original default: $(N, T, T_{\text{burn}}, d, s, k, \rho, \delta, \sigma_x, \sigma_\varepsilon) = (1000, 200, 30, 8, 8, 6, 0.5, 0.5, 0.5, 0.1)$; simplified screening uses $(N, T, d, k) = (300, 120, 6, 4)$.

Each period every agent draws $d_i \sim \text{Bin}(K - m_i, p_{\text{demand}})$ demand slots where m_i is already-active matches and capacity $K=5$. Agent i chooses channel $\in \{\text{:self}, \text{:broker}\}$ by comparing EWMA satisfaction indices. Self-search draws $n_{\text{strangers}}$ random non-neighbours plus all neighbours; the broker offers its current accessible roster ranked by a neural-network prediction over the pair input $(\mathbf{x}_i, \mathbf{x}_j) \in \mathbb{R}^{2d}$. Round-by-round deferred acceptance finalises matches.

2 Instrumentation changes

All edits are additive and backward-compatible: the canonical preflight (`quick_diagnostic.jl` at `seed=42`) produces *byte-identical* numbers pre- and post-instrumentation (outsourcing rate 0.266, R^2 gap 0.426).

2.1 Decoupled fee and cost rates

Legacy model used a single `search_cost_rate` s for both the broker fee and the self-search cost:

$$\phi = c_s = s (q_{\text{cal}} - r).$$

The broker-advantage analysis requires separating these. We add two parameters, with defaults that inherit from s :

$$\phi = s_{\text{broker}} (q_{\text{cal}} - r), \quad s_{\text{broker}} \leftarrow \text{broker_fee_rate (default } s), \quad (3)$$

$$c_s = s_{\text{self}} (q_{\text{cal}} - r), \quad s_{\text{self}} \leftarrow \text{self_search_cost_rate (default } s). \quad (4)$$

Agent satisfaction on the *self* channel updates as $\tilde{q}^{\text{self}} = \sum_{m \in \mathcal{M}_i^{\text{self}}} q_m / d_i - c_s$; on the *broker* channel as $\tilde{q}^{\text{broker}} = (\sum_{m \in \mathcal{M}_i^{\text{brk}}} q_m - \phi |\mathcal{M}_i^{\text{brk, std}}|) / d_i$. Both use *demand slots* d_i as denominator; the self-search cost is charged on every demanded slot whether filled or not, while ϕ is charged only on successful standard broker placements. EWMA update $\text{Sat} \leftarrow (1 - \omega) \text{Sat} + \omega \tilde{q}$ with $\omega = 0.2$.

2.2 Demand, fill, and value counters; algebraic decomposition

We add, per period:

- `self_demand_slots` and `broker_demand_slots` (d^s, d^b), with $d^s + d^b = \text{total_demand}$.
- `self_filled_slots` and `broker_filled_slots` (n^s, n^b), equal to the existing `n_self_matches` and `n_broker_standard`.
- Fill rates $\pi^s = n^s / d^s$, $\pi^b = n^b / d^b$, gap $\Delta_\pi = \pi^b - \pi^s$.
- Gross per-demanded-slot value $\bar{q}^s = (\sum_{m \in \mathcal{M}^s} q_m) / d^s$, similarly $\bar{q}^b$.
- Conditional mean $\bar{q}_{\text{cond}}^s = (\sum q_m) / n^s$, similarly for broker.

Net per-demanded-slot value is

$$\nu^s = \bar{q}^s - c_s, \quad \nu^b = \bar{q}^b - \phi \pi^b, \quad \Delta_\nu = \nu^b - \nu^s, \quad (5)$$

and admits the algebraic decomposition

$$\Delta_\nu = \underbrace{\pi^b (\bar{q}_{\text{cond}}^b - \bar{q}_{\text{cond}}^s)}_{\text{quality-selection}} + \underbrace{(\pi^b - \pi^s) \bar{q}_{\text{cond}}^s}_{\text{fill/access}} + \underbrace{c_s - \pi^b \phi}_{\text{fee/cost}}. \quad (6)$$

We verify this identity to floating-point tolerance on every run (max residual observed: 4.44×10^{-16}). Period-level agent welfare is $W_t^{(a)} = \sum_{m \in \mathcal{M}_t} q_m - c_s d_t^s - \phi n_t^b$; broker revenue $W_t^{(b)} = \phi n_t^b$; both cumulated post-burn.

2.3 Pre-match access/assessment classification

For every accepted brokered match (i, j) we record flags *before* the match-edge and partner-history mutations:

$$\begin{aligned}\text{edge_pre} &= \mathbb{K}\{j \in N_G(i) \text{ before this match}\}, \\ \text{prior_pair_history_pre} &= \mathbb{K}\{\text{partner_count}[i, j] > 0\}.\end{aligned}$$

Matches partition into three disjoint classes:

Class	Definition	Spec interpretation
<code>access_new_edge</code>	$\neg \text{edge_pre}$	broker bridges a structural hole
<code>assessment_reachable_no_prior</code>	$\text{edge_pre} \wedge \neg \text{prior}$	broker ranks a reachable stranger
<code>assessment_prior_partner</code>	prior	broker ranks a known partner

2.4 Ablation modes

An `ablation::Symbol` parameter ($\in \{\text{:none}, \text{:NoBroker}, \text{:NoRegime}, \text{:BlindBroker}, \text{:NoAccessBroker}, \text{:FrozenGraph}, \text{:NoTurnover}, \text{:FreezeBrokerLearning}\}$) gates the model. `:NoBroker` is the paired counterfactual for every treatment run (forces `channel` $\leftarrow$ `:self` for every demander *after* calling `outsourcing_decision` so the RNG state is preserved). `:NoRegime` sets $\delta \leftarrow 0$ at env construction. `:FrozenGraph` suppresses the `add_match_edge!` call in `finalize_accepted_proposal!`. `:NoTurnover` skips entry/exit. `:FreezeBrokerLearning` retains the initial warm-up training but skips `train_broker_nn!` during production periods. `:NoAccessBroker` filters broker candidates to pairs with $j \in N_G(i)$.

BLINDBROKER is the key diagnostic: the broker NN architecture stays at input dimension $2d$, but at every prediction and training call the focal-agent half is zeroed. Concretely,

$$\hat{q}^{\text{blind}}(\mathbf{x}_i, \mathbf{x}_j) = \text{NN}_b(\mathbf{0}_d, \mathbf{x}_j), \quad (7)$$

so the broker retains its cross-agent pooled training data (same width h_b , same history) but has no pair-specific input. Under BLINDBROKER the training uses symmetry-augmented examples $\{(\mathbf{0}, \mathbf{x}_j, q), (\mathbf{0}, \mathbf{x}_i, q)\}$.

`:EqualCapacityBroker` is deferred; the seven implemented ablations cover the decomposition the current analysis needs.

2.5 Holdout metrics: within-agent, pooled, de-meaned, residualized

The existing per-agent holdout samples $n_a = 100$ focal agents and $n_p = 40$ random partners per agent. We report four gap flavours:

- (a) *Within-agent rank/ R^2* : per-agent Spearman ρ and per-agent R^2 on that agent’s holdout, averaged across sampled agents. R^2 denominator is $\sum (q_j - \bar{q}_i)^2$ where $\bar{q}_i$ is the within-agent mean.
- (b) *Pooled rank/ R^2* : Spearman and R^2 on all $n_a n_p$ triples concatenated; global mean replaces $\bar{q}_i$. Pooled- R^2 punishes a predictor that fails to track per-agent shifts.
- (c) *Within-agent de-meaned R^2* : for each sampled i , subtract within-agent means from both truth and prediction before computing R^2 on residuals; average. Formally

$$R_{\text{dm},i}^2 = 1 - \frac{\sum_j ((\hat{q}_{ij} - \bar{\hat{q}}_i) - (q_{ij} - \bar{q}_i))^2}{\sum_j (q_{ij} - \bar{q}_i)^2}.$$

Removes any per-agent offset advantage.

- (d) *Two-way residualized R^2 /rank*: a *fixed* partner sample of $n_p = 40$ is drawn once per period (via a side-car RNG seeded from (seed, t) so the main stream is unperturbed; this preserves the `seed=42` byte-equivalence). All $\min(50, N)$ sampled focal agents are evaluated on the same partner set, giving an estimable (focal $\times$ partner) matrix Q . We subtract the truth’s two-way main effects

$$q_{ij} - \hat{\mu} - \hat{\alpha}_i - \hat{\beta}_j$$

from both truth and predictor, and compute R^2 and Spearman on the *pair-specific residual*. A predictor that captures only additive (row/column) structure scores ≈ 0 ; only pair-conditioned signal shows up.

The residualized metrics are masked to NaN when the pair-specific variance falls below $\max(\sigma_\varepsilon^2 n_p/6, 10^{-8}) \cdot n_p$ and below $10^{-6} \cdot \text{Var}(Q)$ (handles the additively separable $\rho = 1$ case).

2.6 Full/Blind/Agent additive decomposition

For any cell with both a `:none` run (full broker) and a `:BlindBroker` run at the same parameters, we report

$$\text{gap}_{\text{full}} = \underbrace{\text{gap}_{\text{blind}}}_{\text{pooled-data component}} + \underbrace{\text{gap}_{\text{full}} - \text{gap}_{\text{blind}}}_{\text{pair-input component}} \quad (8)$$

separately for within-rank, demeaned- R^2 , residualized-rank, and residualized- R^2 . By construction the "agent-alone gap" in each flavour is zero (it is the reference), so the two components are interpretable in absolute terms.

2.7 Broker fill-failure reason buckets

Unfilled broker-channel slots are attributed to two causes observable from the round-matching workspace:

- `broker_unfilled_pref_exhausted`: demander walked through its complete admissible preference list without holding a final proposal.
- `broker_unfilled_capacity_limited`: demander retained untried preferences but max rounds were reached — typically partners filled by earlier proposers.

The two counters partition `broker_demand_slots - broker_filled_slots` exactly (verified in a smoke test).

3 Diagnostic results

All tables report post-burn means. Gaps are broker-minus-agent except in Table 3 where we split into full-broker and blind-broker. Seed counts are 5 (Stages 1 and 3–4) or 10/arm (Stage 2 pilot, treatment vs. `:NoBroker` control).

3.1 Stage 1: sanity at simplified scale ($N=300, T=120$)

3.2 Stage 2b: full-scale diagnostic ($N=1000, T=200$)

3.3 Stage 2c: Full vs. Blind additive decomposition

Interpretation. At EasyQuality ($\rho = 1$) the DGP (1) is additively separable so *no* pair-specific signal exists; accordingly $\text{rank}_{\text{resid}}$ is near zero for both arms (+0.022 full, −0.006 blind). Yet within-agent rank gap remains +0.235 for BLINDBROKER, **94% of the full-broker value** (+0.249). The broker’s advantage is therefore a data-pooling / sample-size effect: each agent’s

Table 1: Sanity signs. `pi_broker=pi_broker_slot`; "pos" is share of seeds with `rank_gap>0`. Five of seven spec-expected signs hold strictly; `NoRegime` and `EasyQuality` do not show the expected "advantage shrinks", flagged for full-scale retest.

Cell (overrides)	<code>pi_broker</code>	<code>rank_gap</code>	<code>pos</code>	R^2_{gap}	<code>fill_gap</code>	<code>adopt</code>
<code>NoRegime</code> ($\delta=0$)	0.413	+0.096	1.00	+0.447	-0.054	0.78
<code>EasyQuality</code> ($\rho=1$)	0.422	+0.214	1.00	+0.740	-0.050	0.80
<code>HardPairwise</code> ($\delta=.75, \rho=.25$)	0.524	+0.046	1.00	-0.255	-0.080	0.88
<code>HighTransparency</code> (big n_{str}, k)	0.291	+0.048	1.00	+0.196	-0.036	0.74
<code>LowRosterAccess</code> ($\alpha_R=.05$)	0.198	+0.077	1.00	+0.245	-0.091	0.59
<code>HighFee</code> ($\phi\text{-rate} = 0.60$)	0.217	+0.047	0.80	+0.210	-0.071	0.69
<code>FreezeBrokerLearning</code>	0.231	-0.182	0.20	-0.001	-0.035	0.63

Table 2: Full-scale diagnostic. Columns are pooled-data-preserving ratings of the broker-over-agent gap. `EasyQuality` shows a within-agent advantage nearly $2\times$ the Default value, triggering Phase C2c decomposition. $R^2_{\text{residualized}}$ for `EasyQuality` is masked in the patched build (additively separable DGP); here we report the unpatched raw values.

Cell	π_b	<code>rank_within</code>	<code>rank_pooled</code>	R^2_{within}	R^2_{demeaned}	<code>fill_gap</code>
Default ($\delta=.5, \rho=.5$)	0.295	+0.119	+0.124	+0.350	+0.350	-0.051
<code>NoRegime</code> ($\delta=0$)	0.321	+0.127	+0.119	+0.303	+0.386	-0.047
<code>EasyQuality</code> ($\rho=1$)	0.440	+0.249	+0.183	+0.918	+0.570	-0.051
<code>HardPairwise</code> ($\delta=.75, \rho=.25$)	0.363	+0.093	+0.090	-0.015	+0.025	-0.076

neural network trains on ~ 60 partner matches, while the broker pools $\sim 60\,000$ cross-agent observations. At `Default` and `HardPairwise` `BLINDBROKER` *collapses* on residualized rank (-0.60 and -0.62), identifying a genuine pair-conditioned component. `HardPairwise`'s small full-broker lift reflects a *capacity* ceiling: the $h_b=32$ broker NN struggles to fit the discontinuous $\text{sign}(\mathbf{x}_i^\top B \mathbf{x}_j)$ regime gate — planned for a future `EqualCapacityBroker` sensitivity.

3.4 Stage 3: Phase C3 main sweeps at simplified scale ($N=300, T=120, 15$ seeds/arm, treatment + :NoBroker control)

Four 2D sweeps over the base-model parameter space: `delta_rho`, `fee_cost`, `transparency_access`, `learning_turnover`. Each sweep submitted as a slurm job array with one cell per array task, 15 seeds per cell, two arms (treatment and `:NoBroker`). Reducer produces `<sweep>_cell_summaries.csv` and `<sweep>_regime_labels.csv` in `data/summaries/broker_advantage/`.

Regime classifier. For each cell the classifier of `scripts/run_broker_advantage.jl:classify_regime` uses treatment-arm post-burn means and paired ΔW_{ag} :

- $\pi_b > 0.50, \Delta W_{\text{ag}} > 0, \text{rank_gap} > 0 \Rightarrow \text{BROKERDOMINANT}$. Subtype by `access_new_edge_frac_brk`: $> 0.60 \rightarrow \text{Access (BA)}$; $< 0.30 \rightarrow \text{Assessment (BI)}$; else Dual (BD).
- $\pi_b > 0.50, \Delta W_{\text{ag}} < 0 \Rightarrow \text{RENTEXTRACTION (RE)}$.
- $\pi_b < 0.20, \Delta W_{\text{ag}} > 0, \text{rank_gap} > 0 \Rightarrow \text{UNDERADOPTED_INFO (UI)}$.
- $\pi_b < 0.20, \text{rank_gap} > 0, \text{gross_gap} > 0, \Delta W_{\text{ag}} \not\leq 0 \Rightarrow \text{FEEINDUCEDBYPASS (FB)}$.
- $\pi_b < 0.20, \Delta W_{\text{ag}} \not\leq 0, \text{rank_gap} \not\leq 0 \Rightarrow \text{SELFSEARCHDOMINANT (SS)}$.
- Otherwise MIXED (M).

Threshold-robustness maps at $\pi_b=0.30$ and $\pi_b=0.70$ are generated for `delta_rho` and `fee_cost`.

Table 3: Full-broker vs. BLINDBROKER comparison. "Pooled-data component" = blind gap (absolute rank/demeaned- R^2 advantage that BLINDBROKER retains without pair input); "pair-input component" = full – blind (additional lift from seeing $\mathbf{x}_i$). At **EasyQuality** (additively separable DGP) 94% of the within-rank advantage survives blinding — the broker's advantage is overwhelmingly *data pooling*. At **Default** and **HardPairwise** the BlindBroker rank_residualized falls to -0.6 , confirming real pair-conditioned signal.

Cell	arm	π_b	rank _{within}	rank _{resid}	R^2_{within}	R^2_{dm}	fill_gap
Default	full	0.295	+0.119	+0.074	+0.350	+0.350	-0.051
Default	blind	0.204	-0.354	-0.605	-2.690	-0.780	-0.068
EasyQuality	full	0.440	+0.249	+0.022	+0.918	+0.570	-0.051
EasyQuality	blind	0.426	+0.235	-0.006	-0.342	+0.346	-0.048
HardPairwise	full	0.363	+0.093	+0.060	-0.015	+0.025	-0.076
HardPairwise	blind	0.239	-0.581	-0.622	-4.652	-1.703	-0.104

Data integrity caveat on residualized R^2 . At simplified-scale ($N=300, T=120$) the NaN-gate on pair-residualized R^2 fails at $\rho \in \{0.75, 1.00\}$: `r2_gap_residualized_mean` is reported as $+1.34$ to $+1.85$ for $\rho=0.75$ and $+33.06$ (SE 4.48) for *all* four $\rho=1.00$ cells in `delta_rho_cell_summaries.csv`; true $R^2 \in (-\infty, 1]$, so these values exceed range. This is a known gate-calibration issue (`prompts/della_job_guide.md` §15.6) that did fire correctly at full scale (**EasyQuality** NaN in `diagnostic_cell_summaries.csv`), but the gate threshold was miscalibrated for low $N \cdot T$. **We do not cite `r2_gap_residualized` in any claim for cells with $\rho \geq 0.75$ at simplified scale.** Rank-based residualized metric (`rank_gap_residualized_mean`) is bounded in $[-2, 2]$ and remains trustworthy. Also note: at $\rho=1$ the DGP (1) is δ -invariant, so cells `d0.00_r1.00`, `d0.25_r1.00`, `d0.50_r1.00`, `d0.75_r1.00` produce byte-identical summaries.

delta_rho sweep ($4 \times 5 = 20$ cells, $\delta \in \{0, 0.25, 0.5, 0.75\}$, $\rho \in \{0, 0.25, 0.5, 0.75, 1.0\}$). Regime distribution at default thresholds (`delta_rho_regime_labels.csv`): 3 RENTEXTRACTION, 1 BROKERDOMINANT_ASSESSMENT, 16 MIXED. All three RE cells are in the $\rho=0$ column ($\delta \in \{0, 0.25, 0.5\}$); the single BI cell is $(\delta, \rho)=(0.75, 0)$. Cells with the strongest broker adoption are in the $\rho=0$ column (pure pairwise-interaction DGP), where π_b ranges from 0.566 to 0.587.

Table 4: Representative `delta_rho` cells (`data/summaries/broker_advantage/delta_rho_cell_summaries.csv`). All values are post-burn seed means (15 seeds each, treatment arm), except ΔW_{ag} which is the seed-paired treatment–control cumulative welfare. R^2_{within} is the per-agent R^2 gap; “—” marks residualized metrics excluded for $\rho \geq 0.75$ per the caveat above.

Cell (δ, ρ)	π_b (se)	rank_gap (se)	R^2_{within}	net_val_gap	ΔW_{ag} (se)	regime
(0.00, 0.00)	0.582 (0.017)	+0.106 (0.011)	+0.118	-0.232	-69.1 (402.2)	RE
(0.00, 0.50)	0.362 (0.010)	+0.100 (0.012)	+0.374	-0.173	+328.5 (389.6)	M
(0.00, 1.00)	0.414 (0.010)	+0.228 (0.009)	+0.918	-0.078	-836.3 (665.3)	M
(0.25, 0.25)	0.456 (0.017)	+0.090 (0.011)	+0.105	-0.223	-741.9 (357.0)	M
(0.50, 0.00)	0.585 (0.018)	+0.098 (0.011)	+0.043	-0.295	-1278.5 (855.6)	RE
(0.50, 0.50)	0.352 (0.012)	+0.084 (0.009)	+0.302	-0.211	-580.2 (351.8)	M
(0.75, 0.00)	0.566 (0.015)	+0.083 (0.009)	-0.247	-0.289	+1656.8 (845.9)	BI
(0.75, 0.25)	0.492 (0.016)	+0.074 (0.007)	-0.165	-0.253	+364.7 (648.1)	M
(0.75, 0.50)	0.361 (0.014)	+0.069 (0.008)	+0.058	-0.221	+390.5 (597.4)	M
(0.75, 0.75)	0.305 (0.008)	+0.161 (0.009)	+0.765	-0.165	+350.5 (411.5)	M

Observations strictly from the CSV cells cited above:

- **Paired welfare ΔW_{ag} flips sign across the sweep.** Negative (harmful to agents net-of-paired-control) at $\delta=0$ and low δ (e.g. (0.50, 0.50) = -580.2), positive at every $\delta=0.75$ cell

except (0.75, 1.00) = -836.3. The large SEs (up to ~ 860) mean individual-cell sign judgments require more seeds before firm claims.

- **Net value gap is negative in all 20 cells**, ranging from -0.078 (d0.00_r1.00) to -0.295 (d0.50_r0.00). Consistent with the Phase C2b diagnostic finding that the default ϕ makes broker use unprofitable on a per-demanded-slot basis even where prediction is better.
- **rank_gap is positive in 20/20 cells**; share_positive (fraction of seeds with > 0) is 1.0 in 18/20 and 0.933 in 2/20 (d0.00_r0.25, d0.50_r0.25). The broker predicts better for every cell, but R^2_{within} is *negative* in two $\delta=0.75$ cells (d0.75_r0.00: -0.247; d0.75_r0.25: -0.165), echoing the Phase C2c finding that high- δ , low- ρ cells expose a broker-NN capacity shortfall.
- The $\rho=0$ column has the highest π_b (0.566–0.587) and most negative ΔW_{ag} for $\delta \leq 0.5$; the sole BI cell is at the high- δ corner of this column where the paired welfare flips positive.

fee_cost sweep ($5 \times 5 = 25$ cells, $\phi_{\text{rate}} \in \{0, 0.05, 0.15, 0.30, 0.60\}$, $c_{s,\text{rate}} \in \{0, 0.05, 0.15, 0.30, 0.60\}$; δ, ρ held at 0.50). Regime distribution at default thresholds (fee_cost_regime_labels.csv): 3 BROKERDOMINANT_ASSESSMENT, 22 MIXED. All three BI cells are at the *highest* self-search cost ($c_{s,\text{rate}}=0.60$) with low-to-mid broker fee ($\phi_{\text{rate}} \in \{0, 0.05, 0.15\}$), as expected: when self-search is expensive, broker use is dominant.

Residualized R^2 for fee_cost is well-behaved: r2_gap_residualized_mean ranges from 0.192 to 0.296 across all 25 cells (fee_cost_cell_summaries.csv), well within $[0, 1]$. The gate issue is isolated to delta_rho's high- ρ cells.

No FeeInducedBypass (FB) cells detected on this grid. The classifier requires $\pi_b < 0.20$, but π_b ranges from ~ 0.29 to ~ 0.55 even at the highest fee/lowest cost corner. Whether this rules out FB at the simplified scale or merely means the grid did not go high enough on ϕ is an open question; threshold-robustness maps at $\pi_b=0.30$ and $\pi_b=0.70$ (data/figures/broker_advantage/regime_map_fee_cost) will show whether the picture changes at looser bars. **No RentExtraction (RE) cells detected** either. This differs from delta_rho where RE appeared at $\rho=0$; in fee_cost we are at $\delta, \rho=0.5$ so the $\rho=0$ condition that drove delta_rho's RE is not explored here.

transparency_access sweep ($5 \times 4 = 20$ cells, $n_{\text{strangers}} \in \{0, 2, 5, 10, 20\}$, $\alpha_R \in \{0.05, 0.10, 0.20, 0.40\}$; δ, ρ held at 0.50). Regime distribution (transparency_access_regime_labels.csv): **3 BrokerDominantAssessment** (ns0_aR0.20, ns0_aR0.40, ns2_aR0.40), **2 UnderadoptedInfo** (ns5_aR0.05, ns20_aR0.05), **1 RentExtraction** (ns5_aR0.40), **14 Mixed**. This is the sweep where the base-model classifier fires the richest palette. r2_gap_residualized_mean ranges from 0.15 to 0.62 across all 20 cells — in range. Key cells (transparency_access_cell_summaries.csv):

Table 5: Four corner cells of the transparency_access sweep. π_b rises monotonically with α_R (broker roster share) at every $n_{\text{strangers}}$ value in the grid, confirming the expected direction.

Cell (n_{str}, α_R)	π_b (se)	rank_gap	net_val_gap	fill_gap	access_new_edge
(0, 0.05)	0.391 (0.015)	+0.215	-0.204	-0.109	0.319
(0, 0.40)	0.642 (0.009)	+0.203	-0.082	-0.012	0.236
(20, 0.05)	0.181 (0.014)	+0.057	-0.322	-0.080	0.221
(20, 0.40)	0.499 (0.012)	+0.090	-0.141	-0.011	0.222

The $n_{\text{strangers}}=0$ row has the strongest rank_gap and highest broker adoption (up to 0.642 at $\alpha_R=0.40$), because with no self-accessible strangers the broker's roster becomes the primary channel for unknown-partner matches. As $n_{\text{strangers}}$ rises to 20, self-search gains breadth, π_b drops by roughly half, and the rank_gap shrinks accordingly. ΔW_{ag} has large SEs (up to ~ 980),

so sign judgments require caution; the $n_{\text{str}}=0$ column shows positive ΔW_{ag} (72 503 to 76 374 across α_R), which is striking but must be re-read in absolute units with the fee structure noted.

learning_turnover sweep ($4 \times 5 = 20$ cells, $p_{\text{demand}} \in \{0.25, 0.50, 0.75, 0.90\}$, $\eta \in \{0, 0.01, 0.02, 0.05, 0.10\}$; δ, ρ held at **0.50**). Regime distribution (`learning_turnover_regime_labels.csv`): **15 Mixed**, **3 BrokerDominant_Dual** (`pd0.25_eta0.05`, `pd0.25_eta0.10`, `pd0.50_eta0.10`), **1 BrokerDominant_Assessment** (`pd0.25_eta0.02`), **1 RentExtraction** (`pd0.25_eta0.01`). All five non-MIXED cells cluster at low $p_{\text{demand}} \in \{0.25, 0.50\}$ with non-zero turnover $\eta \geq 0.01$; at $p_{\text{demand}} \geq 0.75$ all 10 cells are MIXED.

Table 6: Selected `learning_turnover` cells, all values from `learning_turnover_cell_summaries.csv`. The $\eta=0$ column has *negative* within-agent rank_gap everywhere — no turnover means agents accumulate enough own-model training to outperform a frozen broker on their holdout. Every non-zero η column has positive rank_gap that rises with η .

Cell (p_d, η)	π_b	rank_gap	net_val_gap	ΔW_{ag}	access_frac
(0.25, 0.00)	0.439	−0.018	−0.276	−1785.4	0.072
(0.25, 0.10)	0.926	+0.407	+0.139	+4815.2	0.522
(0.50, 0.00)	0.300	−0.063	−0.244	−2412.0	0.057
(0.50, 0.10)	0.506	+0.302	−0.161	+3649.4	0.534
(0.75, 0.00)	0.211	−0.077	−0.204	+372.8	0.069
(0.75, 0.10)	0.367	+0.261	−0.133	+5890.4	0.542
(0.90, 0.00)	0.196	−0.069	−0.185	+1152.9	0.076
(0.90, 0.10)	0.370	+0.265	−0.105	+6343.9	0.531

Two findings from the CSV worth flagging:

- **The $\eta=0$ column has *negative* rank_gap at every p_{demand} value** (range −0.018 to −0.077; `learning_turnover_cell_summaries.csv` rows `pd*_eta0.00`). At zero turnover, agents accumulate enough own-match history over $T=120$ periods to outperform the broker on their *own* holdout sample. Turnover is the driver of the broker’s rank advantage on this sweep — consistent with Stage 1’s FREEZEBROKERLEARNING sanity cell (−0.182 rank_gap).
- **`pd0.25_eta0.10` is the only cell across all 85 cells of the four simplified-scale sweeps with *positive* net_value_gap** (+0.139). It is also the highest-adoption cell observed ($\pi_b=0.926$) with rank_gap +0.407 and $\Delta W_{\text{ag}}=+4815.2$. This is the clearest candidate for an unambiguously broker-favorable regime at simplified scale.

Cross-sweep regime tally. Across the four simplified-scale sweeps (85 cells total):

- M (Mixed): 67 (79%)
- BI (Assessment-type broker-dominant): 8 — the modal BROKERDOMINANT subtype
- RE (RentExtraction): 5
- BD (Dual-type broker-dominant): 3
- UI (Underadopted_Info): 2
- FB, BA, SS: 0 each

FEEINDUCEDBYPASS, BROKERDOMINANT_ACCESS, and SELFSEARCHDOMINANT do not fire anywhere on the simplified-scale grids at default thresholds. Whether this is a threshold artifact or a real property of the simplified regime will be checked against Stage 4 (§4) and the threshold-robustness maps at $\pi_b \in \{0.30, 0.70\}$ (`regime_map_delta_rho_threshold_broker{030,070}.png`, similarly for `fee_cost`).

3.5 Stage 2 pilot ($N=300, T=120$, 10 seeds, treatment + control)

Table 7: Pilot cells. ΔW_{ag} is paired (treatment – no-broker) agent-welfare cumulative. Note that at $\rho=1$ the δ term drops out of (1); the cells `d0.00_r1.00` and `d0.75_r1.00` are therefore DGP-identical and have identical summaries.

Cell	π_b	rank_gap (CI)	R^2_{gap}	ΔW_{ag} (se)	net_val_gap
$\delta=0.00, \rho=0.00$	0.573	+0.119 [0.10,0.14]	+0.186	+1172.9 (479.8)	–0.208
$\delta=0.00, \rho=1.00$	0.405	+0.222 [0.21,0.24]	+0.866	+812.2 (571.1)	–0.081
$\delta=0.50, \rho=0.50$	0.375	+0.085 [0.06,0.11]	+0.345	+106.2 (398.9)	–0.206
$\delta=0.75, \rho=0.00$	0.603	+0.071 [0.05,0.09]	–0.278	+1581.3 (1187.9)	–0.273
$\delta=0.75, \rho=1.00$	0.405	+0.222 [0.21,0.24]	+0.866	+812.2 (571.1)	–0.081

4 Stage 4: full-scale confirmation ($N=1000, T=200$, 15 seeds, treatment + :NoBroker control)

5 hand-picked cells (from `scripts/run_broker_advantage.jl:full_confirm_cells`): `broker_dominant` ($\delta=0.75, \rho=0.25$); `self_search_dominant` ($\delta=0, \rho=1$); `boundary` ($\delta=0.5, \rho=0.5$); `fee_bypass` ($\delta=0.5, \rho=0.5, \phi_{\text{rate}}=0.60, c_{\text{srate}}=0.05$); `high_info_low_access` ($\delta=0.75, \rho=0.25, \alpha_R=0.05$). Array job 7143835 (split as 10 tasks over (cell, arm) via the new `--arm` CLI option) completed; reducer produced `data/summaries/broker_advantage/full_confirm_cell_summaries.csv`, `full_confirm_regime_labels.csv`, and `full_confirm_paired_summaries.csv` (5 cells, 15 paired welfare deltas each).

Table 8: Stage 4 simplified vs. full comparison. “Simplified” values read from the sweep CSVs at the matching cell: `delta_rho_cell_summaries.csv` for rows 1–3, `fee_cost_cell_summaries.csv` row `bfr0.60_ssc0.05` for `fee_bypass`. `high_info_low_access` has no direct simplified-grid counterpart because $\alpha_R=0.05$ at $\delta=0.75, \rho=0.25$ is off all four 2D grids. All full values from `full_confirm_cell_summaries.csv`. **No regime label flips** between simplified and full at default thresholds; all 5 cells classify as MIXED in both regimes. Magnitudes of ΔW_{ag} amplify substantially at full scale.

Cell	Simpl. π_b	Full π_b	Simpl. ΔW_{ag}	Full ΔW_{ag}	Full rank_gap
<code>broker_dominant</code>	0.492	0.383	+364.7 (648)	+16614.0 (2152)	+0.098
<code>self_search_dominant</code>	0.414	0.432	–836.3 (665)	+1809.9 (1852)	+0.251
<code>boundary</code>	0.352	0.309	–580.2 (352)	+1572.5 (1362)	+0.125
<code>fee_bypass</code>	0.168	0.136	–1780.0 (—)*	–6297.5 (1590)	+0.107
<code>high_info_low_access</code>	—	0.231	—	+20229.6 (2499)	+0.090

* Simplified ΔW_{ag} for `bfr0.60_ssc0.05` read from the `fee_cost` sweep CSV.

All 5 full cells: regime code = M (per `full_confirm_regime_labels.csv`).

Key Stage 4 findings strictly from the CSV cells cited.

- **Magnitude amplification.** At `broker_dominant`, ΔW_{ag} grows from +364.7 (SE 648, not statistically significant at simplified scale) to +16614.0 (SE 2152, clearly > 0) at full scale — a $\sim 45\times$ amplification. Similar amplification in magnitude for the boundary and self-search cells (signs flip but the signal is still within 1–2 SE at simplified scale).
- **Fee-induced bypass signature confirmed.** At `fee_bypass`, full-scale $\pi_b=0.136$ is below the 0.20 threshold for FEEINDUCEDBYPASS, rank_gap is +0.107, and ΔW_{ag} is –6297.5 (SE 1590). The paired welfare is *negative and statistically significant*: in this fee regime, agents

would be strictly better off under the `:NoBroker` counterfactual. The regime classifier still labels this MIXED rather than FB because FB additionally requires `gross_gap > 0` (broker’s *gross* quality per demanded slot exceeds self’s); with π_b this low, the broker’s per-demanded-slot numerator shrinks and the comparator fails. **We flag this as an “open question” per runbook §11.2 Stage-4 rule:** a fee-induced-bypass semantically-correct cell that does not fire the strict classifier, suggesting the `gross_gap` condition should be relaxed or removed for this regime.

- **high_info_low_access** ($\delta=0.75, \rho=0.25, \alpha_R=0.05$) has the largest positive ΔW_{ag} across all full-confirm cells (+20229.6 with SE 2499). This cell could not be directly matched in any simplified 2D sweep. Interpretation: when the regime signal is strong (high δ) but self-search access is narrow (low α_R), paired welfare moves sharply in the broker’s favor.
- **self_search_dominant** ($\rho=1$) does *not* actually show dominant self-search at full scale: $\pi_b=0.432$, `rank_gap` = + 0.251. The broker’s pooled-data advantage carries over to this cell (consistent with the Phase C2c finding). The cell’s name reflects the intent (zero regime, zero pair-interaction DGP); the observation is that pooled data still delivers a sizeable rank advantage even where pair-conditioned signal is zero.

5 Summary of findings

Claims in this section cite a specific CSV row, figure, or identity residual. “Pending” marks items awaiting completion of still-running sweeps or Stage 4.

(i) **Pooled-data vs. pair-input decomposition.** At the full-scale diagnostic (`diagnostic_cell_summaries.` Phase C2c blind comparison, Table 3):

- In the **Default** cell ($\delta=0.5, \rho=0.5$), BLINDBROKER (pooled data, no pair input) scores `rank_residualized` = -0.605 while Full scores $+0.074$ (Table 3). The pair-input component (Full – Blind, eq. (8)) is $+0.679$ in residualized-rank units — pair-conditioned signal is the dominant driver here.
- In **EasyQuality** ($\rho=1$, additively separable DGP), `rank_residualized` is near zero for both arms ($+0.022$ Full, -0.006 Blind) as theory predicts, yet within-agent `rank_gap` is $+0.249$ Full vs. $+0.235$ Blind — Blind retains 94% of the within-agent rank advantage. The broker’s apparent edge here is data pooling (cross-agent $\sim 60\,000$ observations vs. ~ 60 per agent), not information asymmetry.
- In **HardPairwise** ($\delta=0.75, \rho=0.25$), Blind collapses to `rank_residualized` = -0.622 while Full stays at $+0.060$ — pair input is again essential; the small Full value reflects the $h_b=32$ broker NN’s capacity shortfall on the discontinuous regime gate.

These full-scale decompositions are the reliable reference; simplified-scale residualized- R^2 in `delta_rho` is gate-broken at high ρ (§3.4) and cannot be used for analogous claims on the sweep grid. Residualized-*rank*, bounded, remains usable at simplified scale.

(ii) **Fee-induced bypass and rent-extraction.**

- From `delta_rho_regime_labels.csv`: 3 cells classify as RENTEXTRACTION (RE): `d0.00_r0.00`, `d0.25_r0.00`, `d0.50_r0.00`. All are in the $\rho=0$ column with $\pi_b \geq 0.58$ and paired $\Delta W_{\text{ag}} < 0$ (Table 4). At default fee/cost rates, when the DGP is pure pairwise-interaction (no additive-quality cue for agents), agents outsource heavily but are paired worse off in cumulative welfare.

- One more RE cell appears in `transparency_access_regime_labels.csv` at `ns5_aR0.40`, and one in `learning_turnover_regime_labels.csv` at `pd0.25_eta0.01`. Total RE across the four sweeps: 5/85 cells.
- FEEINDUCEDBYPASS (FB) requires $\pi_b < 0.20$ with positive `rank_gap`. **No FB cells appear on any of the four sweeps** at default thresholds. In `fee_cost`, π_b ranges 0.29–0.55 even at the high- ϕ /low- c_s corner `bfr0.60_ssc0.00` — the simplified-scale grid does not drive adoption low enough to hit the $\pi_b < 0.20$ FB condition. Threshold-robustness maps at $\pi_b=0.30$ and $\pi_b=0.70$ (`regime_map_fee_cost_threshold_broker{030,070}.png`) are produced but do not change the bypass-vs-dominant decision for this grid. Whether FB emerges at full scale is tested in §4.
- **Net value gap is negative in 84 of 85 cells across all four sweeps.** The single exception is `pd0.25_eta0.10` in the `learning_turnover` sweep, where `net_value_gap_mean` = +0.139 with $\pi_b=0.926$, `rank_gap` = +0.407, ΔW_{ag} = +4815.2 (`learning_turnover_cell_summaries.csv`, Table 6). This is the clearest broker-beneficial regime found at simplified scale: high turnover at low demand frequency gives the broker a sustained information advantage and fills most demanded slots.

(iii) **Regime distribution and the “Mixed” majority.** 67/85 cells across all four sweeps classify as MIXED (M) under default thresholds ($\pi_b=0.50$, access threshold 0.60). The remaining 18 cells are: 8 BI, 5 RE, 3 BD, 2 UI. The thresholds are conservative; the alternative maps at $\pi_b \in \{0.30, 0.70\}$ are produced for `delta_rho` and `fee_cost`.

(iv-bis) **Turnover is the driver of broker rank advantage.** In `learning_turnover`, every cell with $\eta=0$ has *negative* `rank_gap` (range -0.018 to -0.077 across $p_{\text{demand}} \in \{0.25, 0.5, 0.75, 0.9\}$, Table 6). Every cell with $\eta>0$ has positive `rank_gap` that rises monotonically with η . This aligns with the Stage 1 FREEZEBROKERLEARNING sanity result (`rank_gap` = -0.182), and suggests the broker’s informational edge is sustained by entry/exit replacing agents before they accumulate enough own-match history to outperform the broker’s pooled model.

(iv) Cross-sweep findings.

- *fee_cost*: No FEEINDUCEDBYPASS (FB) or RENTEXTRACTION (RE) cells on the 5×5 grid at $\delta, \rho=0.5$. BROKERDOMINANT_ASSESSMENT (BI) emerges only at the highest self-search cost $c_{s,\text{rate}}=0.60$ with low-to-mid fee (`fee_cost_regime_labels.csv` rows `bfr0.00_ssc0.60`, `bfr0.05_ssc0.60`, `bfr0.15_ssc0.60`). The `fee_sensitivity_curves.png` figure, read alongside, shows whether adoption or net value drops monotonically with ϕ_{rate} .
- *transparency_access*: Four distinct regime codes appear on a single sweep (3 BI, 2 UI, 1 RE, 14 M) — this sweep is the richest diagnostic palette at simplified scale. UNDERADOPTED_INFO (UI: $\pi_b < 0.20$ with positive `rank_gap` and ΔW_{ag}) appears at high $n_{\text{strangers}}$ and low α_R — exactly where the broker’s informational advantage is real but its access is too narrow for agents to adopt. RE appears uniquely at $n_{\text{str}}=5$, $\alpha_R=0.40$. π_b is monotonic in α_R at every fixed $n_{\text{strangers}}$ (Table 5), confirming that roster size is a primary lever for adoption.
- *learning_turnover*: pending (job 7142787 complete, reducer 7143900 running).
- *Phase C4 ablations (completed, 24/24 tasks)*: 3 anchors $\times$ 8 modes, simplified scale, 15 seeds per cell, 1 arm (treatment). Full/Blind decomposition written to `data/summaries/broker_advantage/ablat`

At **baseline** and **broker_dominant**, the pair-input component is large and positive, and the pooled-data component is *negative* (BlindBroker is worse than a per-agent model). At

Table 9: Ablation full/blind decomposition at simplified scale (`ablations_decomposition.csv`). “Pooled data” is the BLINDBROKER arm’s gap (zero pair input, same pooled training). “Pair input” is Full – Blind. Residualized- R^2 values at `self_search` anchor ($\rho=1$) exceed $(-\infty, 1]$ as flagged in §3.4; *we do not cite them here*.

Anchor	Metric	Full	Blind	Pooled-data	Pair-input
baseline ($\delta=0.5, \rho=0.5$)	rank_within	+0.079	−0.416	−0.416	+0.495
baseline	rank_residualized	+0.037	−0.641	−0.641	+0.678
baseline	R^2_{demeaned}	+0.258	−0.903	−0.903	+1.162
broker_dominant ($\delta=0.75, \rho=0.25$)	rank_within	+0.058	−0.596	−0.596	+0.654
broker_dominant	rank_residualized	+0.037	−0.660	−0.660	+0.697
broker_dominant	R^2_{demeaned}	−0.204	−1.838	−1.838	+1.634
self_search ($\delta=0, \rho=1$)	rank_within	+0.238	+0.236	+0.236	+0.003
self_search	rank_residualized	+0.057	−0.051	−0.051	+0.109
self_search	R^2_{demeaned}	+0.643	+0.506	+0.506	+0.138

`self_search` (additively separable, $\rho=1$), pair-input component is essentially zero (+0.003 within-rank), and the entire +0.238 full-rank advantage is pooled-data — directly replicating the Phase C2c full-scale conclusion at simplified scale with 15 seeds per arm.

Regime distribution across the 24 ablation cells (`ablations_regime_labels.csv`): 21 MIXED, 3 SELFSEARCHDOMINANT (`baseline_FreezeBrokerLearning`, `baseline_NoBroker`, `broker_dominant_NoBroker`). The `NoBroker` arms correctly classify as SS since $\pi_b=0$ by construction; `baseline_FreezeBrokerLearning` collapsing to SS confirms the Stage 1 sanity finding (−0.182 rank_gap) at the main-anchor scale.

(v) **Stage-4 regime flips.** *Pending.* Will be read from `data/summaries/broker_advantage/full_confirm_s` after job completion and reported in Table 8.

6 Findings so far and open items

1. Sanity criterion “rank_gap shrinks under NoRegime / EasyQuality” is a spec expectation based on equal-data asymptotics. At realistic finite (N, T) the broker’s pooled-data advantage dominates any loss of pair-conditioned structure. The instrumentation is sound — the preflight baseline is byte-identical to the pre-instrumentation run on `seed=42`.
2. **Net value gap Δ_ν is negative in every cell examined** (−0.07 to −0.30). At default ϕ , broker use is not net-beneficial to agents on a per-demanded-slot basis even where prediction is better. This motivates the `fee_cost` and `transparency_access` sweeps.
3. The BLINDBROKER-Full pairwise decomposition (8) cleanly separates pooled-data from pair-input advantages and will be read alongside the regime map for every main-sweep cell at the anchor rows.
4. Pilot criterion (2) — “rank_gap near zero in no-regime / pure-quality anchor” — does not hold strictly at simplified scale; 4/5 criteria pass. Main sweeps launched regardless, conditional on Stage 4 full-scale confirmation.
5. **Known caveat.** Residualized R^2 gate is miscalibrated at simplified scale for $\rho \geq 0.75$; use `rank_gap_residualized` there instead (see §3.4). Full-scale residualized R^2 is correctly NaN-gated (`diagnostic_cell_summaries.csv` EasyQuality row).

Reproducibility

All code is in `TransientBrokerage.jl` at commit `KK-testing-branch` (HEAD at time of writing). New scripts: `scripts/run_broker_advantage.jl` (CLI with `--stage/--sweep/--cell-idx/--reduce`) and `sysimage/build_sysimage.jl` (PackageCompiler sysimage for $\sim 25\times$ TTFX reduction on array tasks). Running notes with seeds, wallclocks, and cluster topology live at `data/summaries/broker_advantage`. The session handoff document is `prompts/della_job_guide.md`.

Phase panels — delta_rho sweep (post-burn means, n_seeds=15)




Figure 1: `delta_rho` sweep phase panels (ρ on the x-axis, δ on the y-axis; treatment arm only, 15 seeds per cell). Panel (1,1) π_b ; (1,2) paired ΔW_{ag} ; (1,3) `net_value_gap`; (2,1) within-agent `rank_gap`; (2,2) within-agent R^2 gap; (2,3) fill-rate gap; (3,1) `access_new_edge` fraction; (3,2) `assessment_prior_partner` fraction; (3,3) regime code. All values are post-burn seed means from `data/summaries/broker_advantage/delta_rho_cell_summaries.csv` and `delta_rho_regime_labels.csv`.




Figure 2: `delta_rho` regime map at default thresholds ($\pi_b > 0.5$ for broker-dominant; $\text{access} > 0.60 \rightarrow \text{BA}$; $< 0.30 \rightarrow \text{BI}$; else BD). Short codes from `regime_map` legend: RE = RENTEXTRACTION, BI = BROKERDOMINANT_ASSESSMENT, M = MIXED. 3 RE cells at $\rho=0$ with $\delta \in \{0, 0.25, 0.5\}$; the single BI cell is at $(\delta, \rho)=(0.75, 0)$; 16 Mixed cells fill the rest.

Phase panels — fee_cost sweep (post-burn means, n_seeds=15)




Figure 3: fee_cost sweep phase panels (broker_fee_rate on x, self_search_cost_rate on y; $\delta, \rho=0.5$ held). Panel layout matches Figure 1. All values from data/summaries/broker_advantage/fee_cost_cell_summaries.csv; regime codes from fee_cost_regime_labels.csv.

Phase panels — transparency_access sweep (post-burn means, n_seeds=15)




Figure 4: transparency_access sweep phase panels ($n_{\text{strangers}}$ on x, α_R on y; $\delta, \rho=0.5$). Layout matches Figure 1. π_b rises monotonically in α_R at every $n_{\text{strangers}}$ (cf. Table 5); four distinct regime codes appear on this sweep (3 BI, 2 UI, 1 RE, 14 M per transparency_access_regime_labels.csv).

Phase panels — learning_turnover sweep (post-burn means, n_seeds=15)




Figure 5: learning_turnover sweep phase panels (p_{demand} on x, η on y; $\delta, \rho=0.5$). Layout matches Figure 1. The $\eta=0$ row has *negative* rank_gap at every p_{demand} ; non-MIXED regimes cluster at low p_{demand} with $\eta \geq 0.01$. Values from learning_turnover_cell_summaries.csv and learning_turnover_regime_labels.csv.

Full-confirm comparison (N=1000/T=200 vs N=300/T=120)




Figure 6: Stage 4 full-confirm comparison across the 5 hand-picked cells at $N=1000, T=200$, 15 seeds/arm (data/summaries/broker_advantage/full_confirm_cell_summaries.csv). Panels: π_b , net_value_gap , rank_gap , R^2_gap . All 5 cells classify as MIXED at default thresholds; the underlying ΔW_{ag} magnitudes amplify at full scale relative to the simplified-sweep counterparts (Table 8).

Fee sensitivity — curves over broker_fee_rate per self_search_cost_rate




Figure 7: Fee sensitivity: π_b , **net_value_gap**, **rank_gap**, and ΔW_{ag} (or net value fallback) as functions of **broker_fee_rate**, one curve per **self_search_cost_rate** level (**fee_cost_cell_summaries.csv**). The absence of a large drop in π_b at the highest fees is the figure-level evidence against a FEEINDUCEDBYPASS regime at this grid's default $(\delta, \rho)=(0.5, 0.5)$.




Figure 8: Representative cells across sweeps, temporal overlay (data/sims/broker_advantage/<sweep>/cell__id__seed*__arm_treatment.jld2). Each panel shows ensemble means across the 15 seeds; burn-in period marked by the vertical dashed line. Panels: π_b , net_value_gap, rank_gap, access_new_edge_frac_brk, betweenness, welfare_agents_period.




Figure 9: Ablation bars by anchor (π_b , rank_gap, R^2 _gap, net_value_gap, access_new_edge_frac_brk) across the 8 ablation modes (none, NoBroker, NoRegime, BlindBroker, NoAccessBroker, FrozenGraph, NoTurnover, FreezeBrokerLearning). Post-burn seed means $\pm$ SE from data/summaries/broker_advantage/ablations_cell_summaries.csv. BlindBroker collapses rank_gap at baseline and broker_dominant but not at self_search (cf. Table 9).